2020-2021学年秋季学期

数据科学导论

The Introduction of Data Science

授课团队：沙瀛　周川

助　　教：梁棋

数据科学导论

The Introduction of Data Science

[第　章] 总结与展望

授课教师：周川

授课时间：2020年12月25日

主要章节

1. 引言
2. 数据科学生命周期
3. 相关与因果-批判性思维
4. 数据预处理
5. 数据分析与建模
6. 数据可视化
7. 项目实施与沟通
8. 分类
9. 聚类
10. 回归
11. 关联规则分析
12. 异常检测
13. 数据降维
14. 时间序列分析

第8章 分类

- 决策树算法
- 朴素贝叶斯分类器
- 最近邻分类器
- Logistics回归
- 提升方法（集成学习）

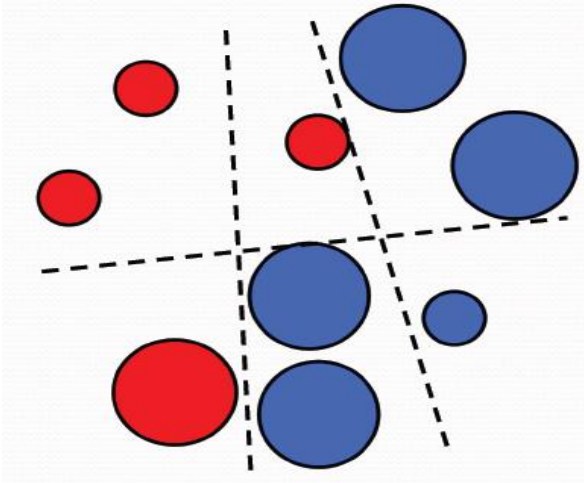

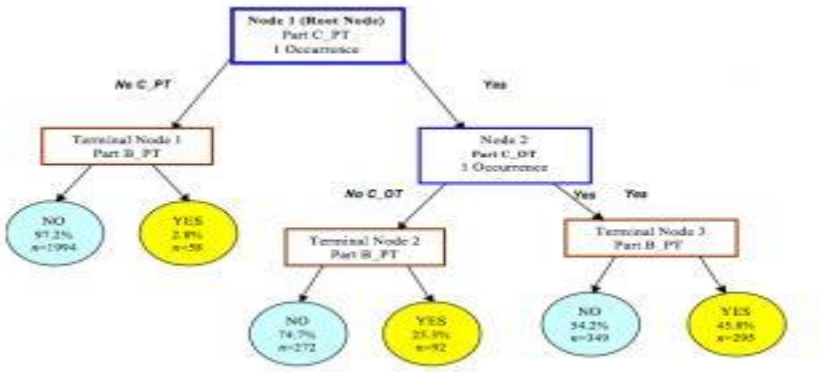

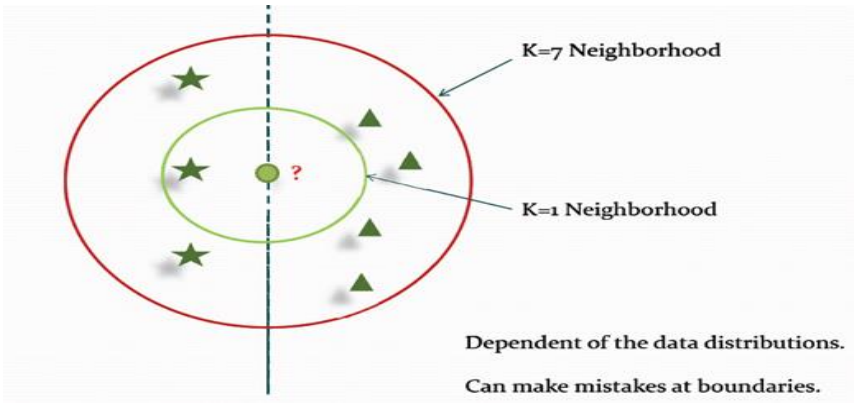

第9章 聚类

- 接近性度量/差异性度量
- 顺序聚类算法
- 划分聚类算法（K-means聚类算法）
- 层次聚类算法（ AGNES 、 DIANA ）
- 密度聚类算法（ DBSCAN ）

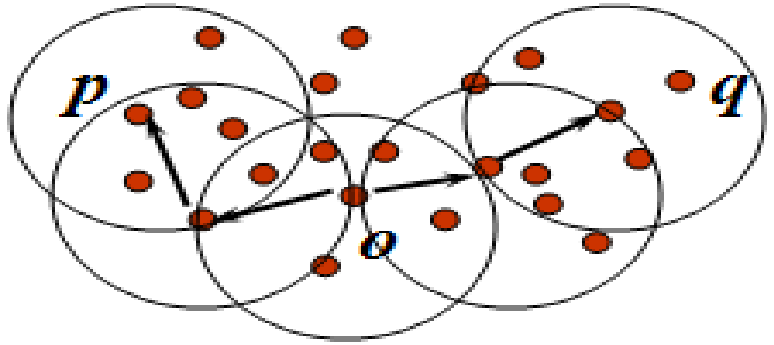

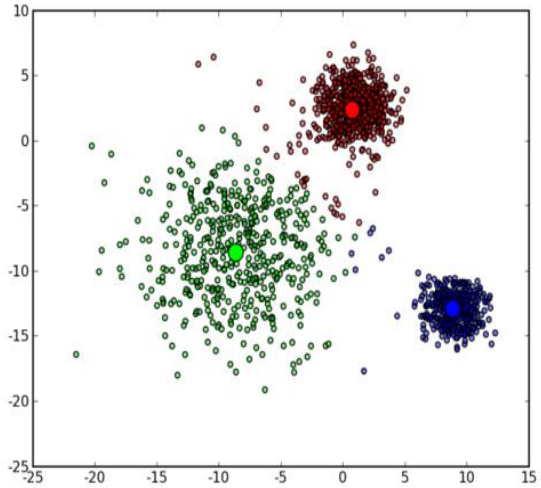

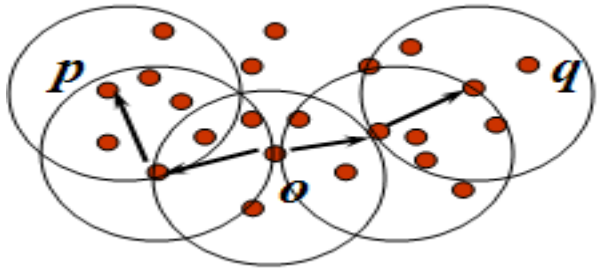

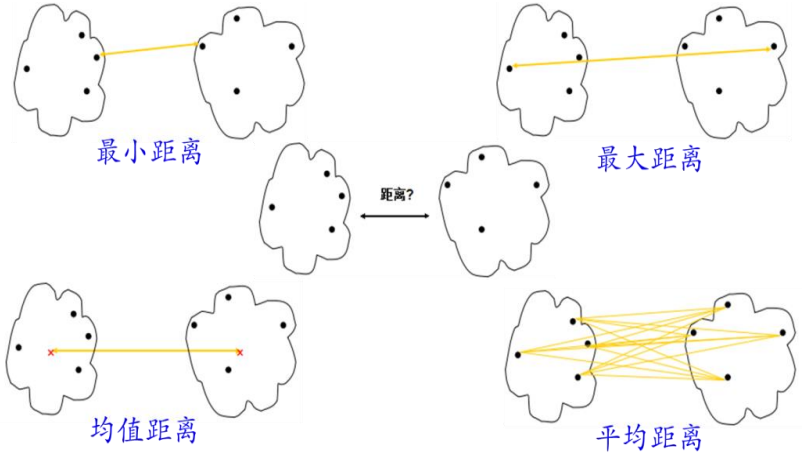

第10章 回归

- 一元线性回归
- 多元线性回归
- 非线性回归
- 附录：逐步回归、岭回归、Lasso回归

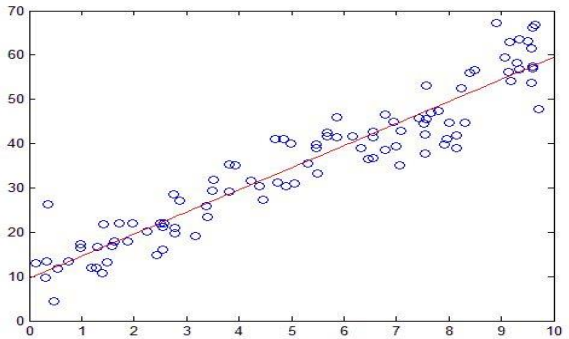

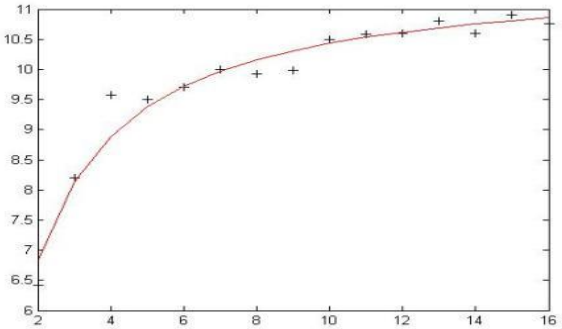

$$\hat{\beta} = (X^{\tau}X)^{-1}X^{\tau}Y$$

$$\begin{cases} y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p + \varepsilon \\ \varepsilon \sim N(0, \sigma^2) \end{cases}$$

第11章 关联规则分析

- 频繁相集/关联规则
- Apriori算法
- PCY算法
- FP-Growth算法
- 序列模式挖掘

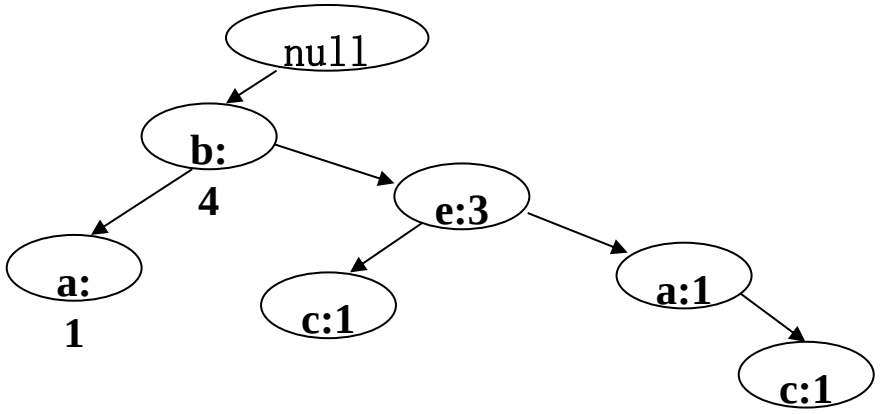

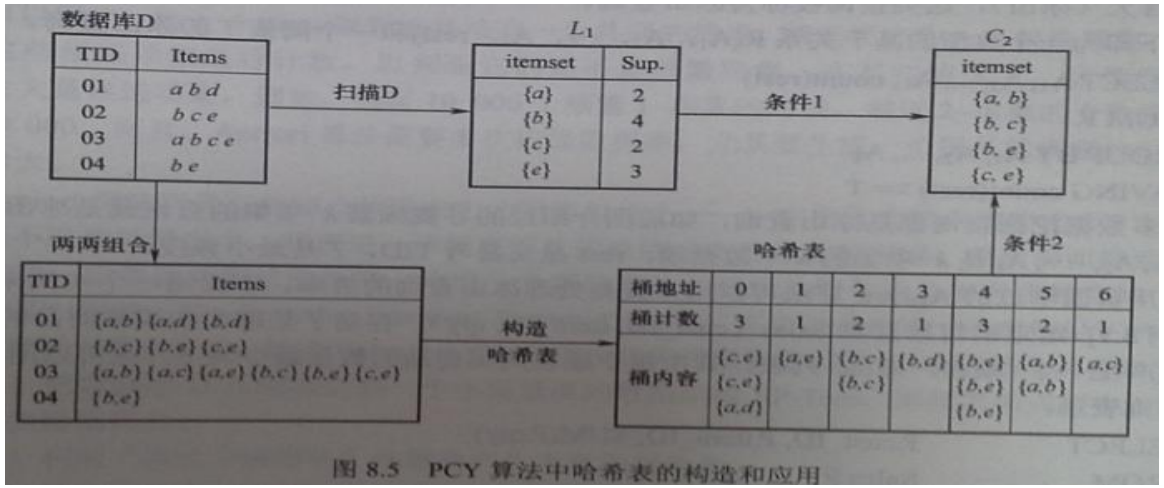

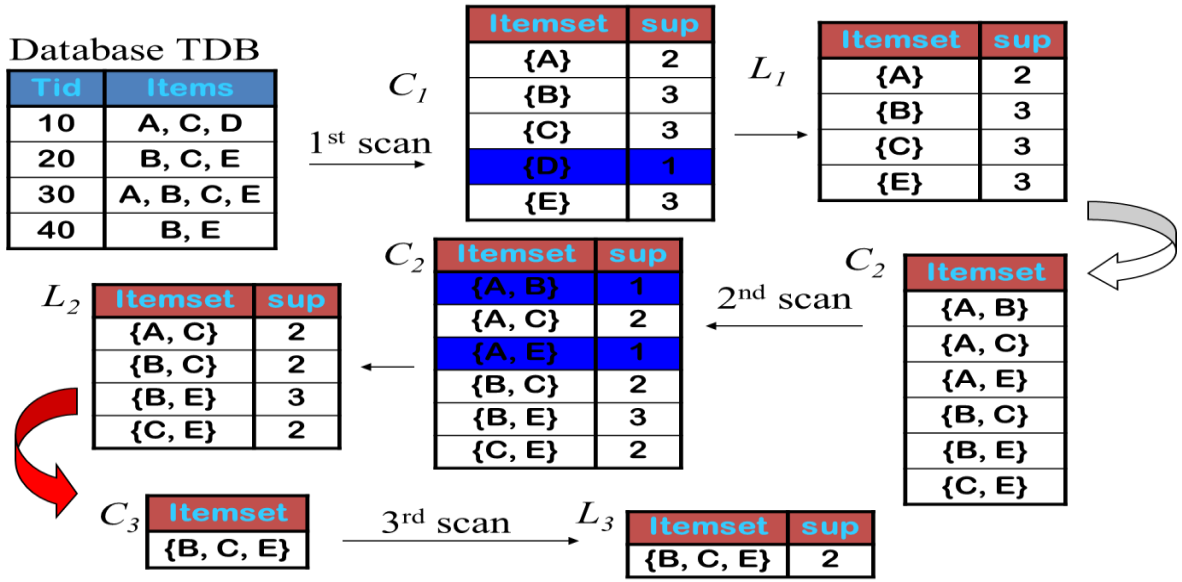

第12章 异常检测

- 定义与主要挑战
- 基于图形的方法
  - 箱型图
- 基于统计的方法
- 分类和聚类
- 基于距离和基于密度的方法
  - 基于K-means的方法
  - 局部异常因子（LOF）方法

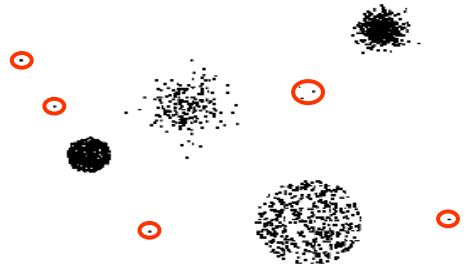

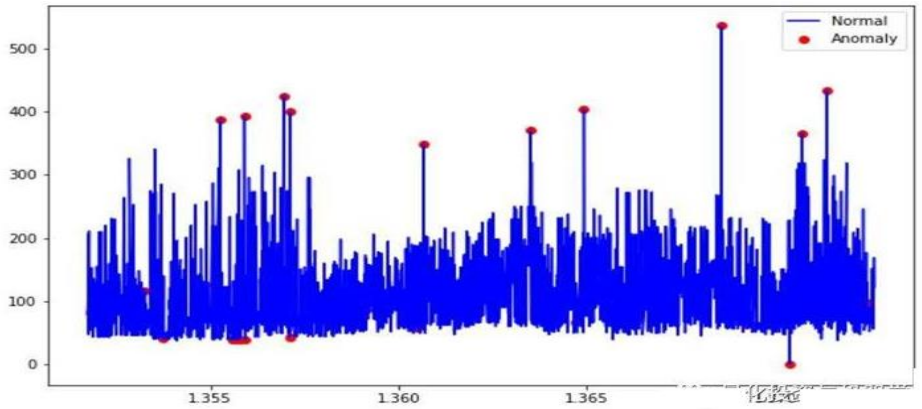

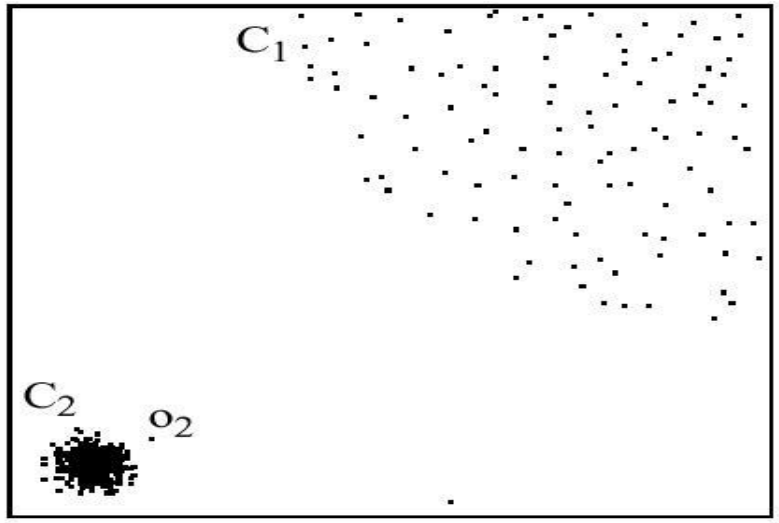

第13章 数据降维

- 主成分分析
  - 几何意义、总体主成分、样本主成分
- 因子分析
  - 因子载荷、因子旋转、因子得分
- SVD分解
- 低维嵌入

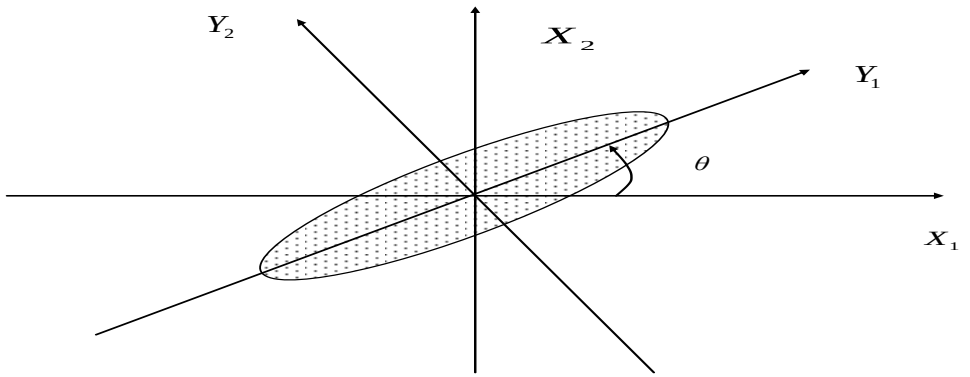

$$\boldsymbol{\Sigma} = \mathbf{AA}' + \mathbf{D} = \mathbf{U}\begin{bmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & & & \lambda_p \end{bmatrix}\mathbf{U}'$$

$$\boldsymbol{A}_{m\times n} = \boldsymbol{U}_{m\times m}\boldsymbol{\Sigma}_{m\times n}\boldsymbol{V}^{\mathrm{T}}_{n\times n}$$

$$\varepsilon(\boldsymbol{W}) = \sum_{i=1}^{n}\left|\boldsymbol{x}_i - \sum_{j=1}^{k} w_{ij}\boldsymbol{x}_{i\_j}\right|^2 \qquad \Phi(\boldsymbol{W}) = \sum_{i=1}^{n}\left|\boldsymbol{y}_i - \sum_{j=1}^{k} w_{ij}\boldsymbol{y}_{i\_j}\right|^2$$

第14章 时间序列分析

- 时间序列分析的预处理
  - 平稳性检验
  - 随机性检验
- 平稳时间序列分析
  - ARMA模型
  - 平稳序列建模与学习
  - 序列预测
- 非平稳时间序列分析
  - 序列分解
  - ARIMA模型
  - 残差自回归模型
  - 指数平滑预测模型

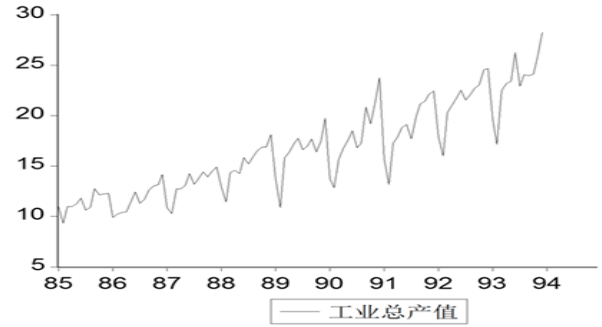

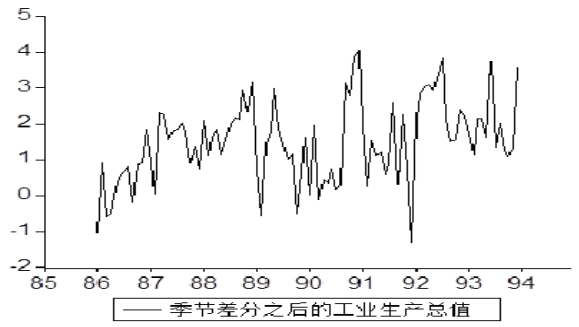

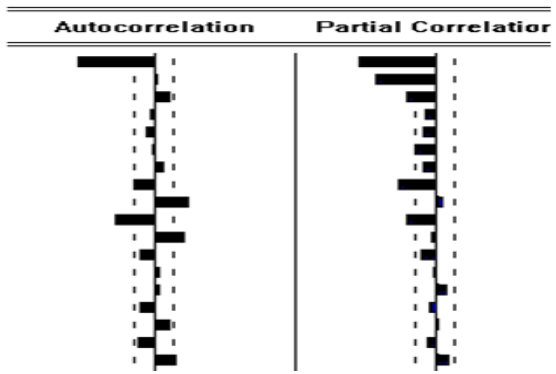

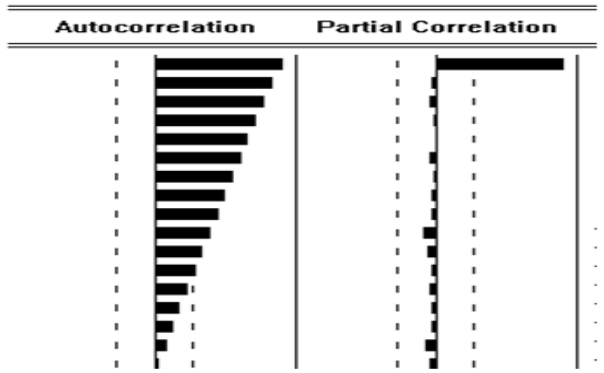

What is Data Science?

Extraction of knowledge from large volumes of data that are structured or unstructured.

It is a continuation of the fields **data mining** and **predictive analytics**

My Definition for Data Science

The application of **data centric**, **computational**, and **inferential thinking** to

understand the world — **Science**

&

solve problems — **Engineering**

➢ *Data science is fundamentally interdisciplinary*

Data Science

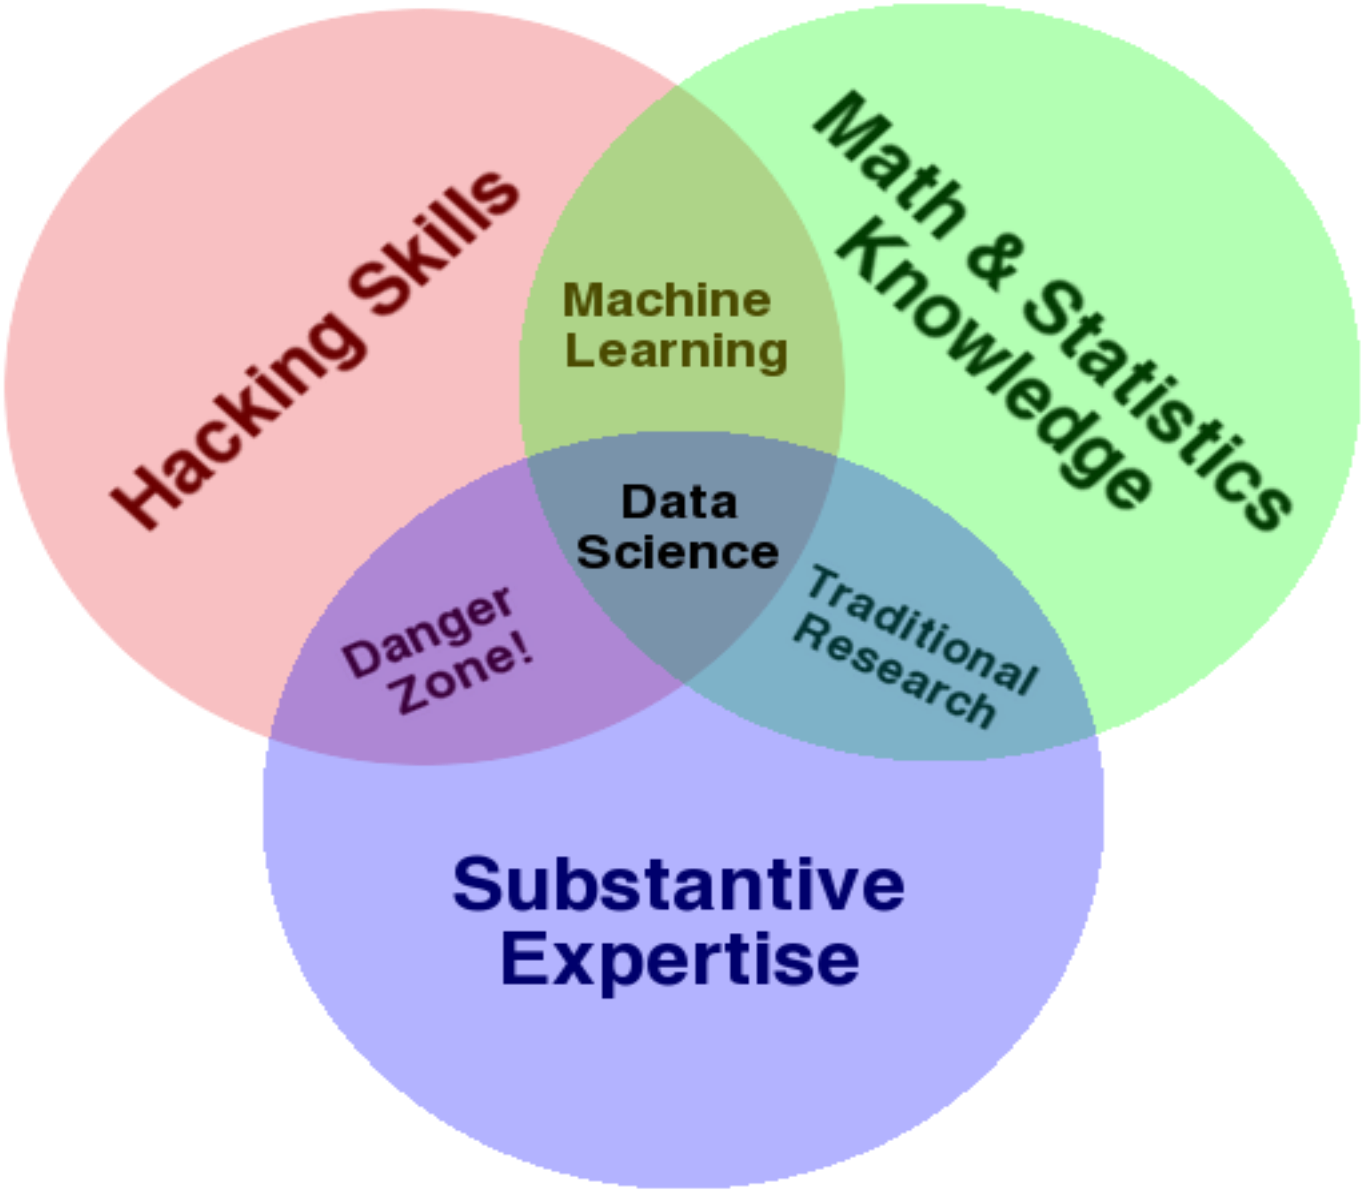

Data Science Tools

python™

ANACONDA®
Powered by Continuum Analytics

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$

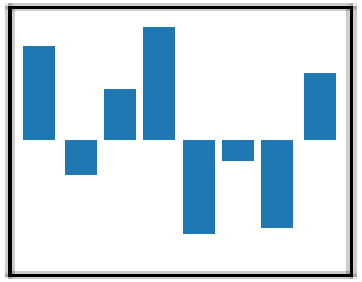

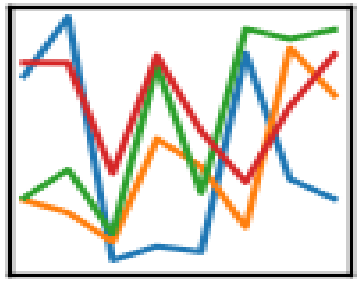

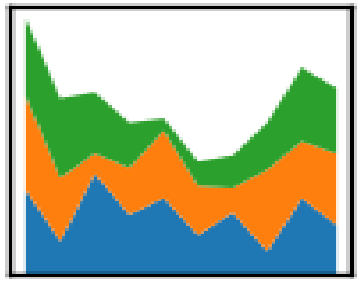

DATA SCIENCE LIFECYCLE: AN ALTERNATE VIEW

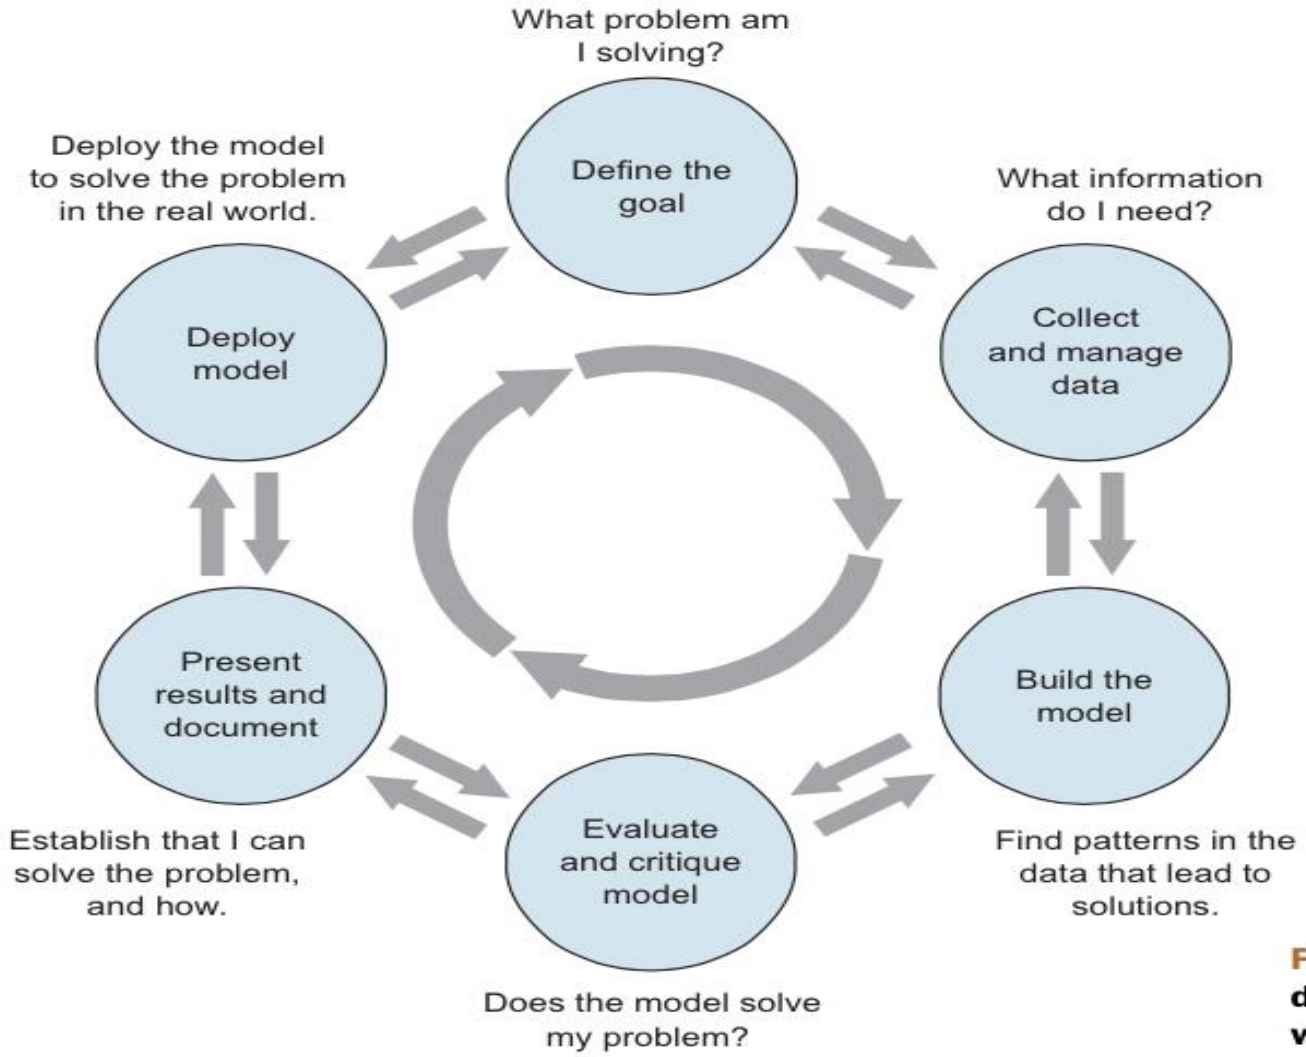

Figure 1.1 The lifecycle of a data science project: loops within loops

DATA SCIENCE LIFECYCLE

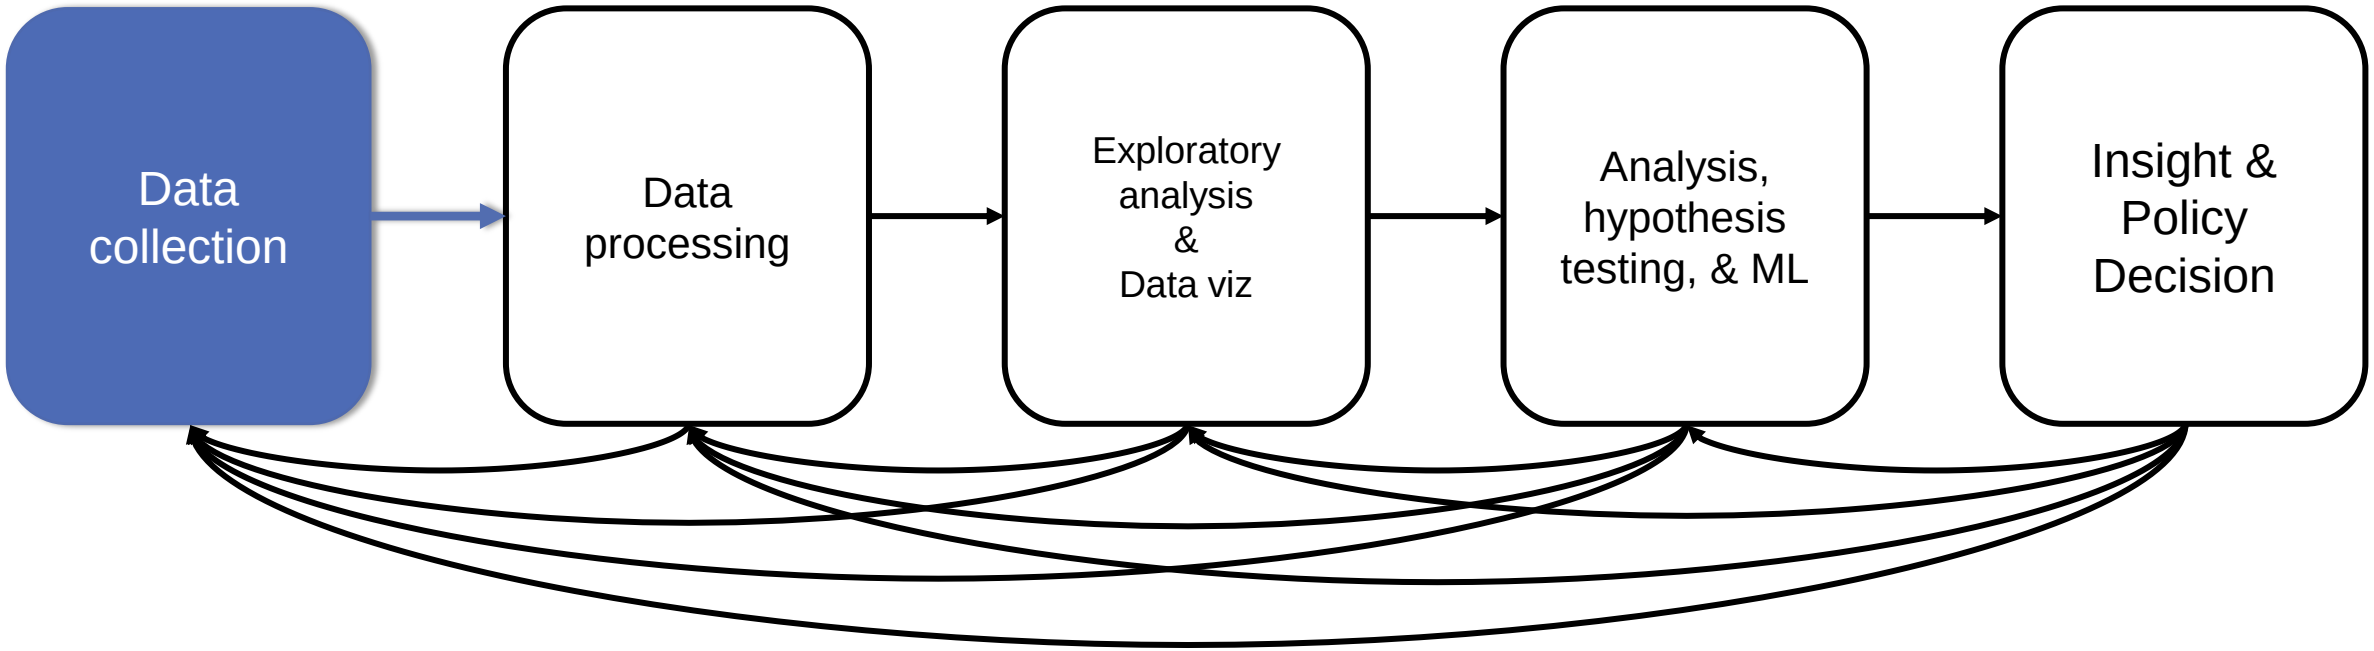

采集的目标：快、准、全

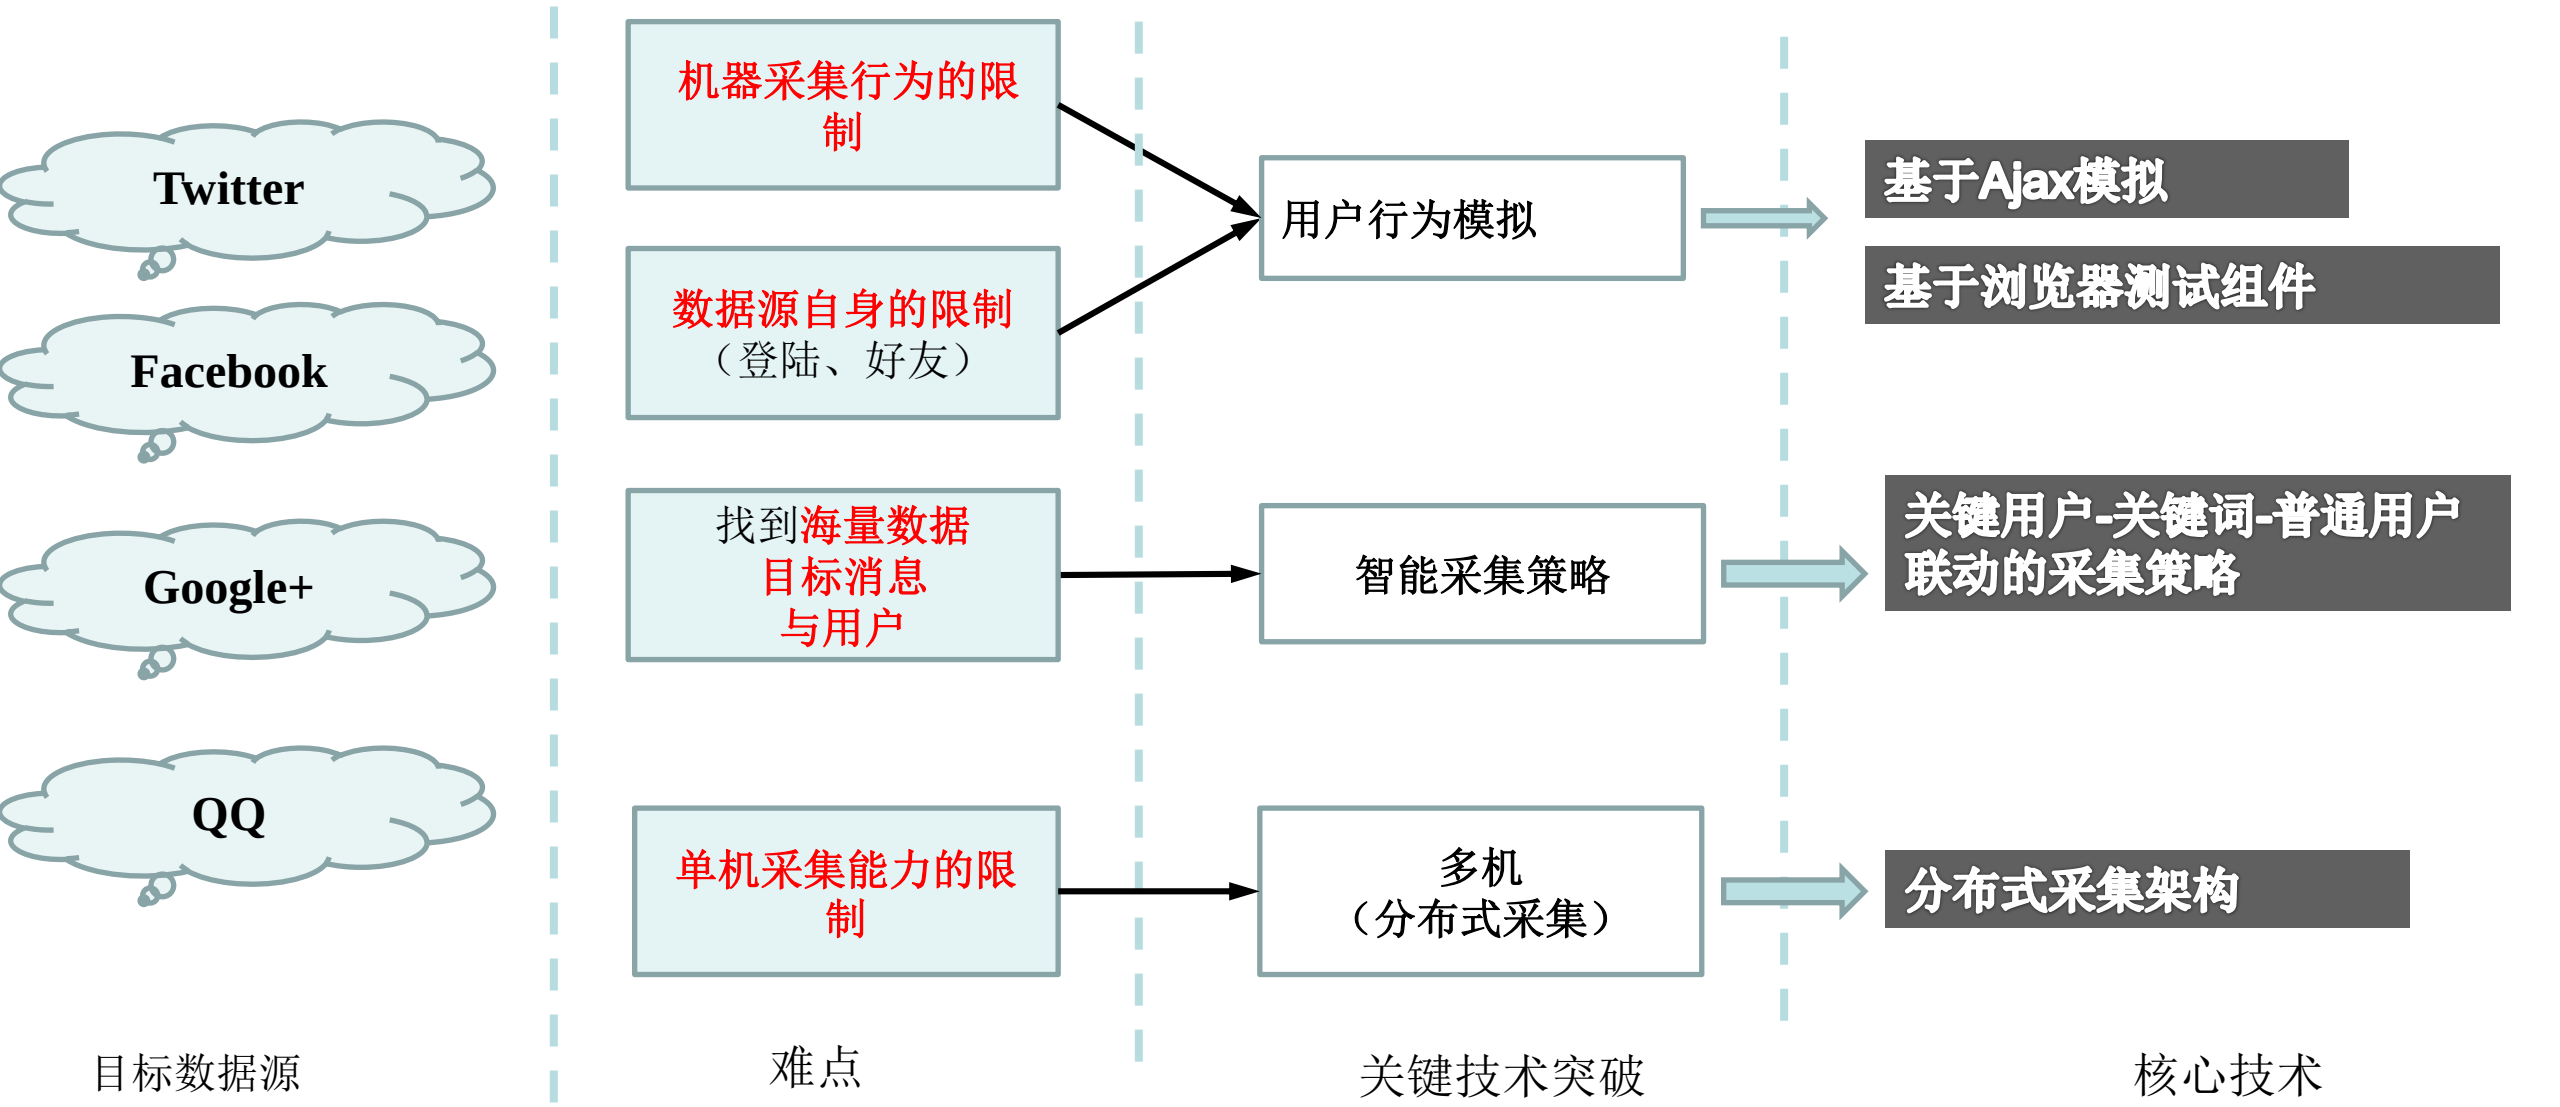
Twitter
Facebook
Google+
QQ
机器采集行为的限制
数据源自身的限制
（登陆、好友）
找到海量数据
目标消息
与用户
单机采集能力的限制
用户行为模拟
智能采集策略
多机
（分布式采集）
基于Ajax模拟
基于浏览器测试组件
关键用户-关键词-普通用户联动的采集策略
分布式采集架构
目标数据源
难点
关键技术突破
核心技术

Data Cleanup

Real data is messay, often needs to cleaned up before useful.

- Bad forms
- Missing Data
- Useless Variables
- Wrong Data

Data Cleanup (Contd.)

- Transform variables (date formats, String to int)
- Create derived variables
  - Derive county from IP
  - Age from ID card number
- Normalize strings
  - Different spelling and nicknames (William -> Bill)
- Feature value rescaling
- Enrich

Data Exploration

Understand, and get a feel for what is expected (models => densities, constraints) and unexpected/ residuals (errors, outliers)

- think what this is data about? domain, background, how it is collected, what each fields mean and range of values.
- head, tail, count, all descriptives (Mean, Max, median, percentiles ..) - Five number Summary. Min. 1st Qu. Median Mean 3rd Qu. Max.
- run a bunch of count/group-by statements to gauge if it's corrupt

Data Exploration (Contd.)

- Plot - take random sample and explore (scatter plot)
  - e.g. Draw scatter plot or Trellis Plot
- Find Dependencies between fields
  - Calculate Correlation
  - Dimensionality reduction
  - Cluster and look visualize clusters
- Look at frequency distribution of each field and try to find a known distribution if possible.

Data Exploration (Contd.)

Feature Engineering

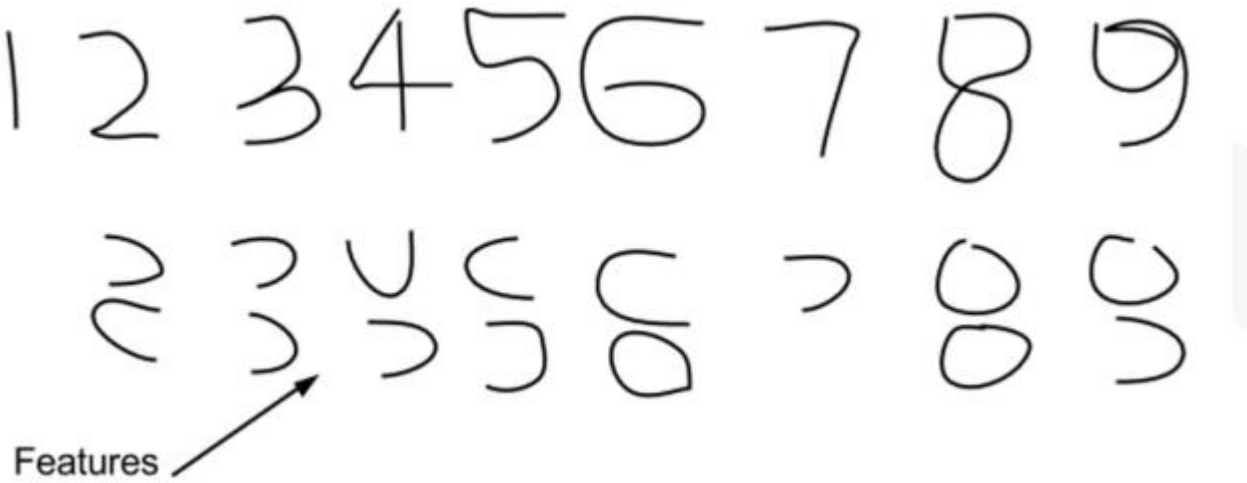

- Feature engineering is the art of finding feature that leads simplest decision algorithm. (Good features allow a simple model to beat a complex model.)
- Best features may be a subset, or a combination, or transformed version of the features.

How to do Feature Engineering?

- Manually pick by domain experts and trial and error.
- Search the possible combinations by training and combining subsets (e.g. Random Forest)
- Use statistical concepts like correlation and information criteria
- Reduce the features to a low dimension space using techniques like PCA.
- Automatic Feature Learning though Deep Learning

Analysis

- Goal of analysis is to extract knowledge
- This knowledge usually come in one of the two forms
  - KPI (Key Performance Indicators)
    - Describe key measurement for what is being measured. (e.g. revenue per year, profit margin, revenue for sqft in retail, revenue per employer)
  - Models to describe or predict the data
    - e.g. Machine Learning models or Statistical models

4 Analysis types by time to decision

- Hindsight (what happened?)
  - Done using Batch Analytics like MapReduce
- Oversight (what is happening?)
  - Done using Realtime Analytics technologies like CEP
- Insight (why things happening?)
  - Done with Data Mining and Unsupervised learning algorithms like Clustering
- Foresight (what will happen?)
  - Done by building models using Machine learning or one of other techniques

Data Analytics Tools : Apache Spark

- Unifies ***batch, interactive, streaming*** workloads
- Easy to build sophisticated applications
  - Support iterative, graph-parallel algorithms
  - Powerful APIs in Scala, Python, Java, R

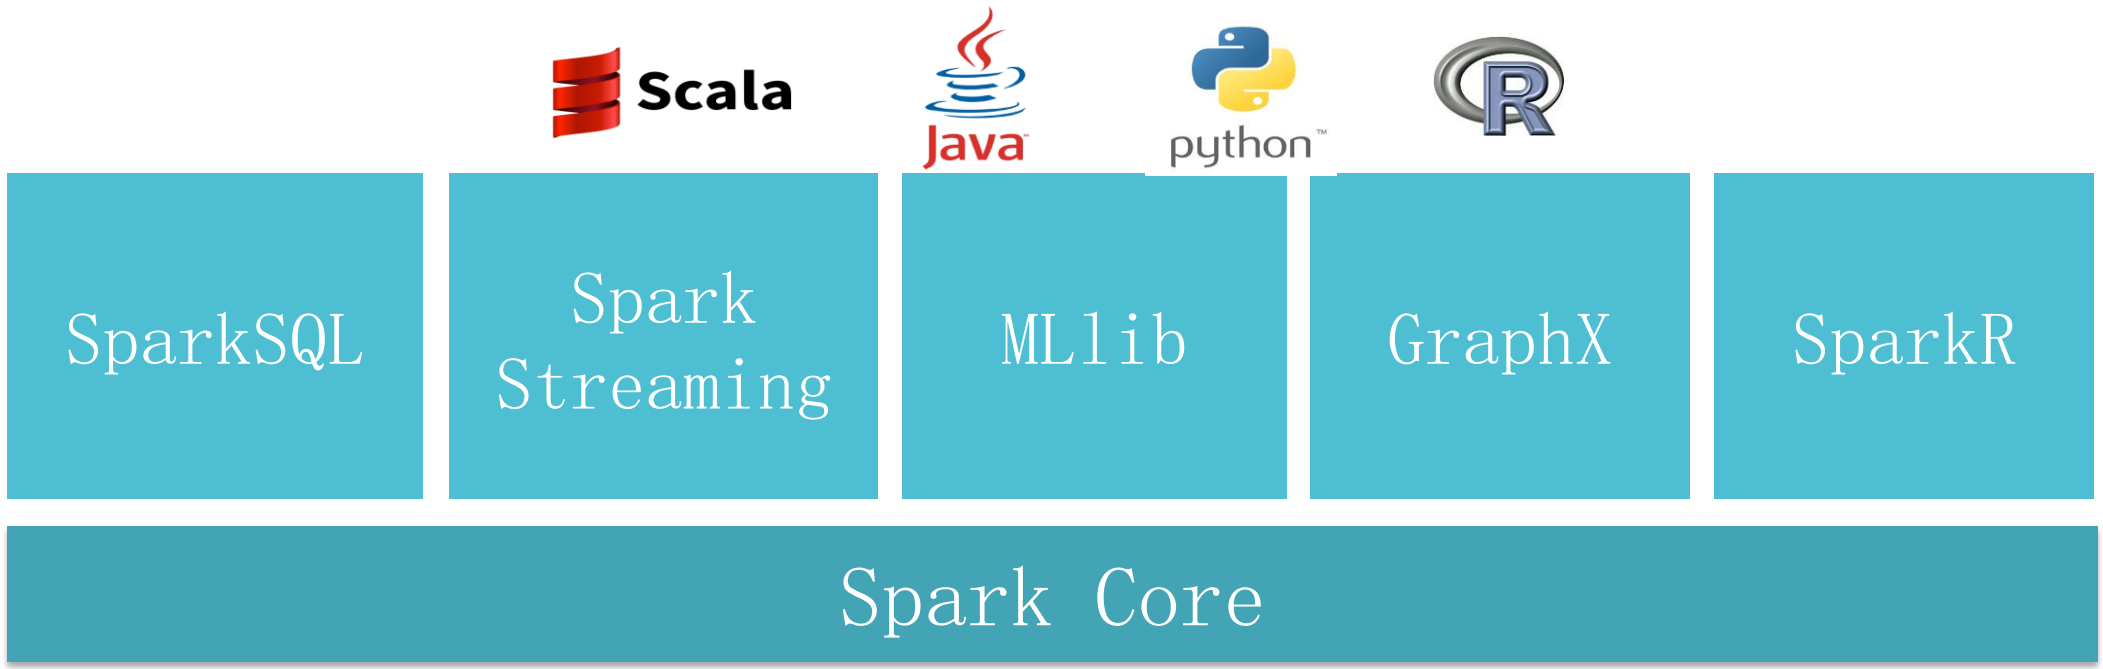

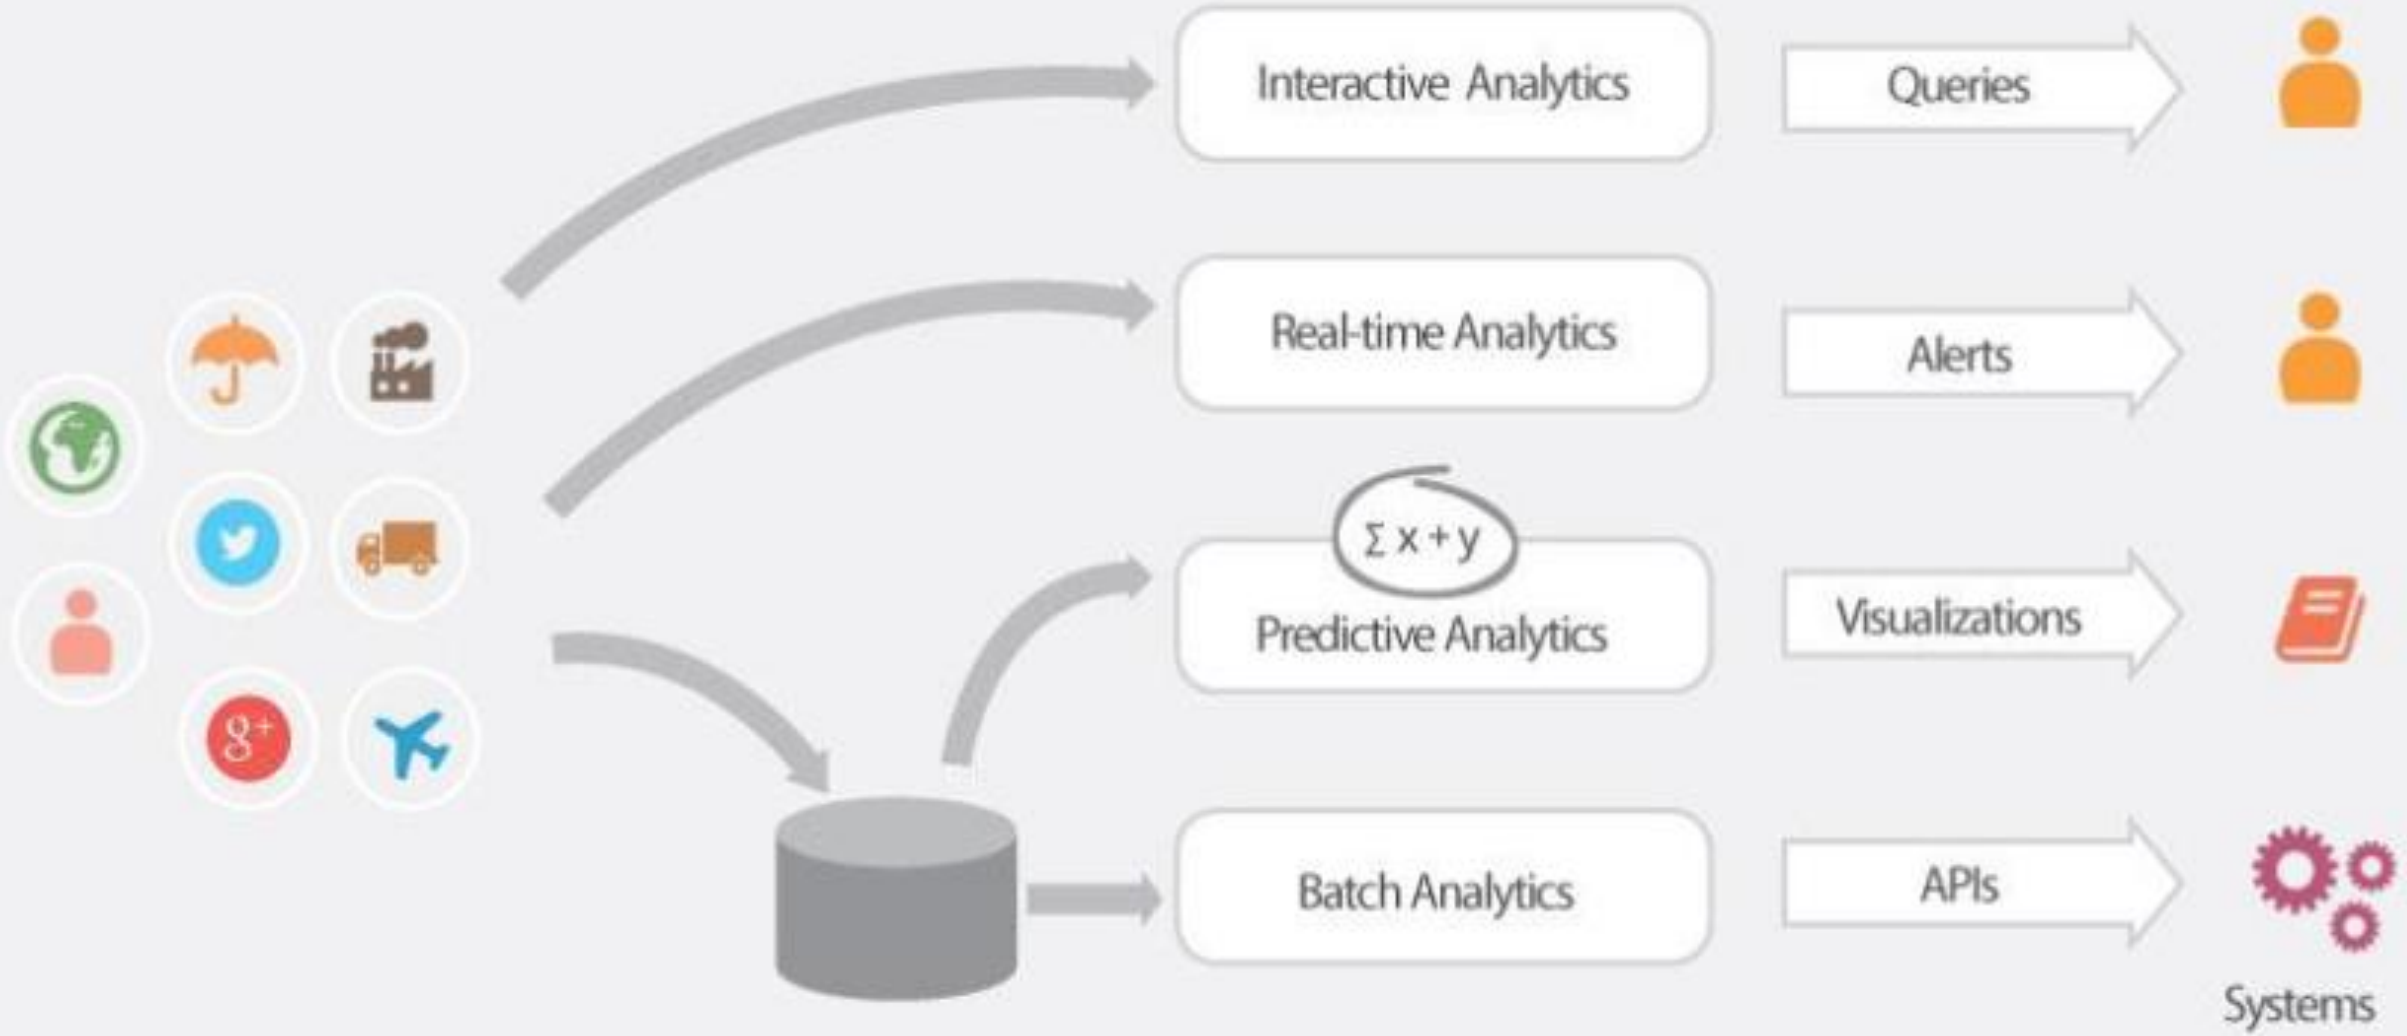

Collect Data → Analyze & Make Decisions → Communicate

Inconvenient Truth About Data Science

- Data is **never clean**
- You will spend **most of your time** cleaning and preparing data.
- **95%** of tasks do not require deep learning
- In **90%** of cases generalized linear regression will do the trick
- Big Data is just a **tool**.
- You should embrace the **Bayesian** approach.
- No one cares **how you did it**.
- Academia and business are **two different** worlds.
- **Presentation is key** – be a master of Power Point.
- All models are false, but **some** are **useful**.
- There is no fully automated Data Science. You need to **get your hands dirty**.

Data Science is extracting knowledge by analyzing data

Hindsight

Oversight

Insight

Foresight

谢谢大家！